

```
In [51]: import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv('vibrations.csv')

print("First 5 rows:")
print(df.head())

print("\nColumn names:")
print(df.columns)

print("\nDataset shape:")
print(df.shape)

signal = df['mean'].values

plt.figure(figsize=(10,4))
plt.plot(signal[:2000], color='hotpink')
plt.title('Mean Vibration Feature Over Time')
plt.xlabel('Window Index')
plt.ylabel('Mean Amplitude')
plt.show()
```

First 5 rows:

| | max | min | mean | sd | rms | skewness | kurtosis | \ |
|---|---------|----------|----------|----------|----------|-----------|-----------|---|
| 0 | 0.35986 | -0.41890 | 0.017840 | 0.122746 | 0.124006 | -0.118571 | -0.042219 | |
| 1 | 0.46772 | -0.36111 | 0.022255 | 0.132488 | 0.134312 | 0.174699 | -0.081548 | |
| 2 | 0.46855 | -0.43809 | 0.020470 | 0.149651 | 0.151008 | 0.040339 | -0.274069 | |
| 3 | 0.58475 | -0.54303 | 0.020960 | 0.157067 | 0.158422 | -0.023266 | 0.134692 | |
| 4 | 0.44685 | -0.57891 | 0.022167 | 0.138189 | 0.139922 | -0.081534 | 0.402783 | |

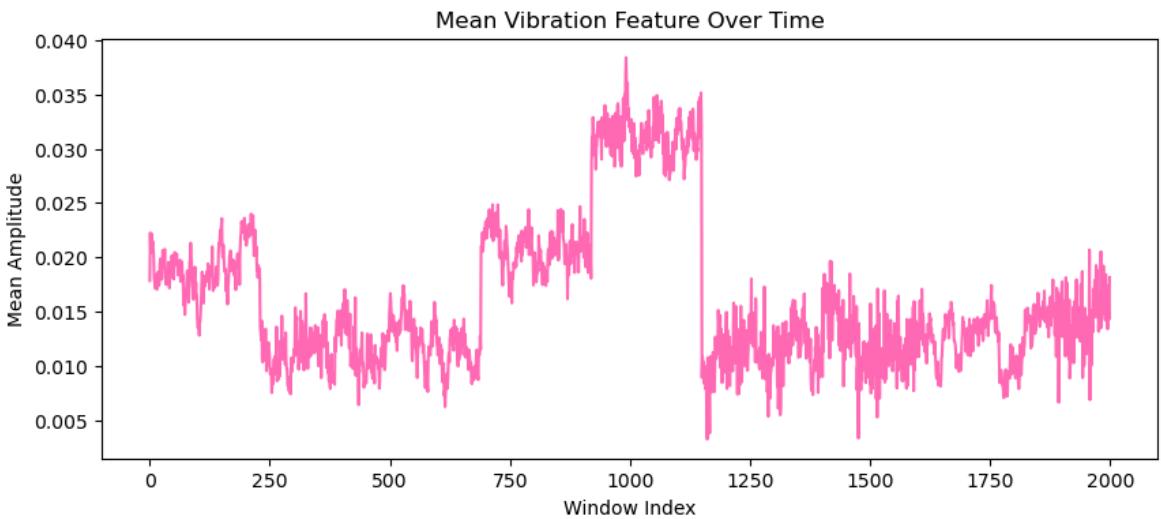
| | crest | form | fault |
|---|----------|----------|------------|
| 0 | 2.901946 | 6.950855 | Ball_007_1 |
| 1 | 3.482334 | 6.035202 | Ball_007_1 |
| 2 | 3.102819 | 7.376926 | Ball_007_1 |
| 3 | 3.691097 | 7.558387 | Ball_007_1 |
| 4 | 3.193561 | 6.312085 | Ball_007_1 |

Column names:

```
Index(['max', 'min', 'mean', 'sd', 'rms', 'skewness', 'kurtosis', 'crest',
       'form', 'fault'],
      dtype='object')
```

Dataset shape:

```
(2300, 10)
```



```
In [36]: print("Fault label distribution:")
print(df['fault'].value_counts())
```

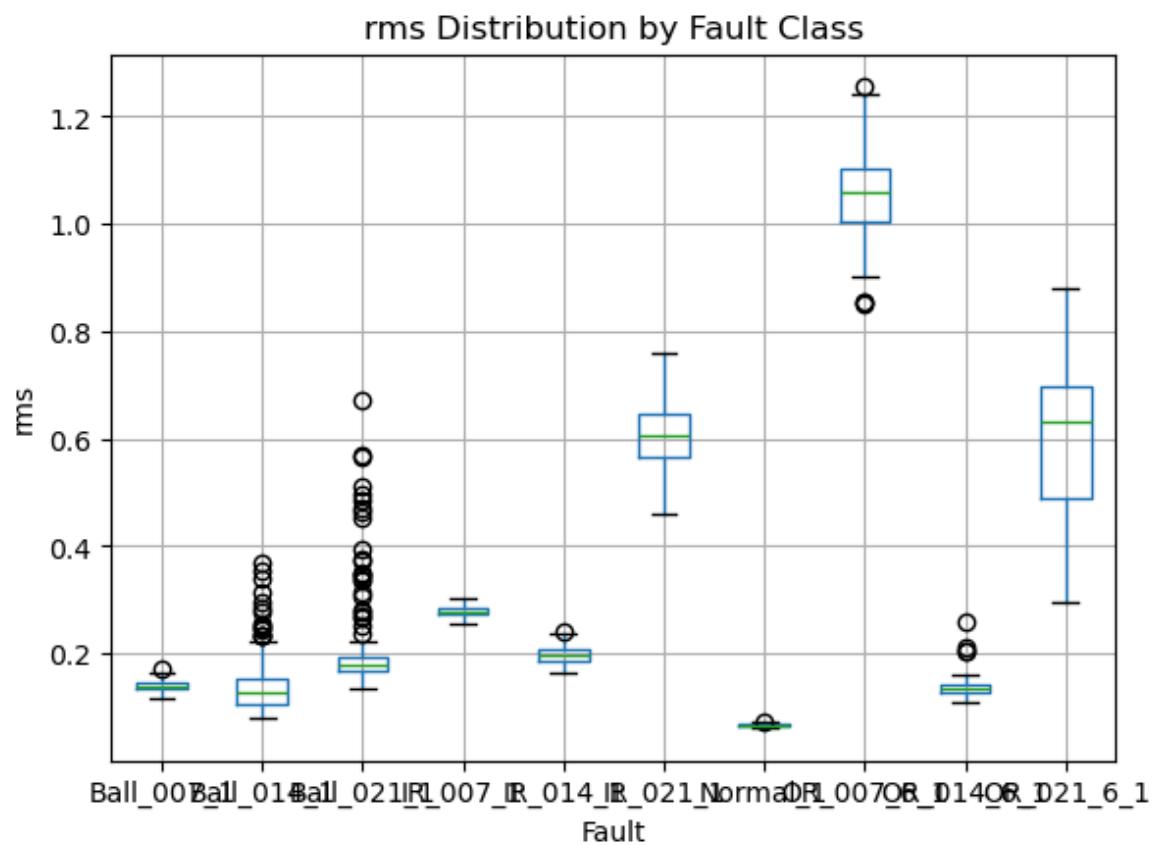
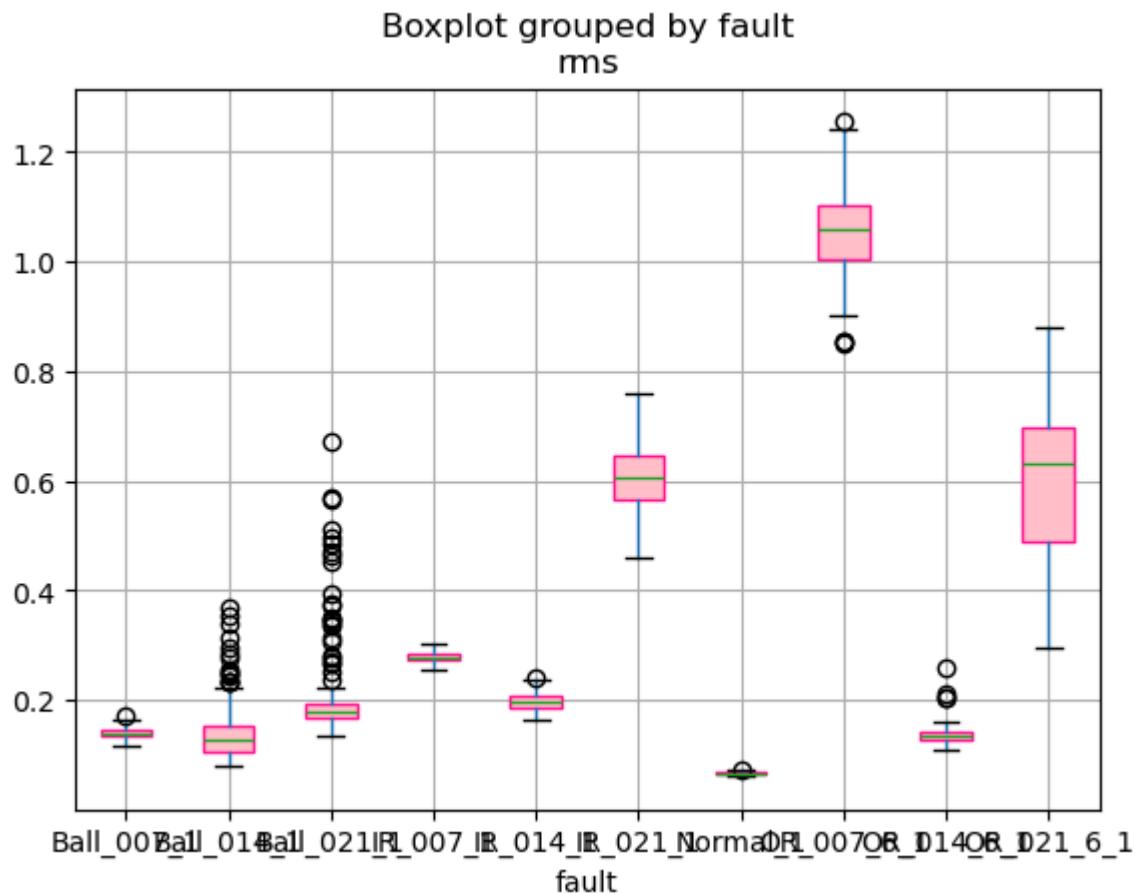
```
Fault label distribution:
fault
Ball_007_1    230
Ball_014_1    230
Ball_021_1    230
IR_007_1      230
IR_014_1      230
IR_021_1      230
OR_007_6_1    230
OR_014_6_1    230
OR_021_6_1    230
Normal_1       230
Name: count, dtype: int64
```

```
In [52]: X = df.drop(columns=['fault'])
y = df['fault']
X.describe()

features = ['rms', 'kurtosis', 'sd']

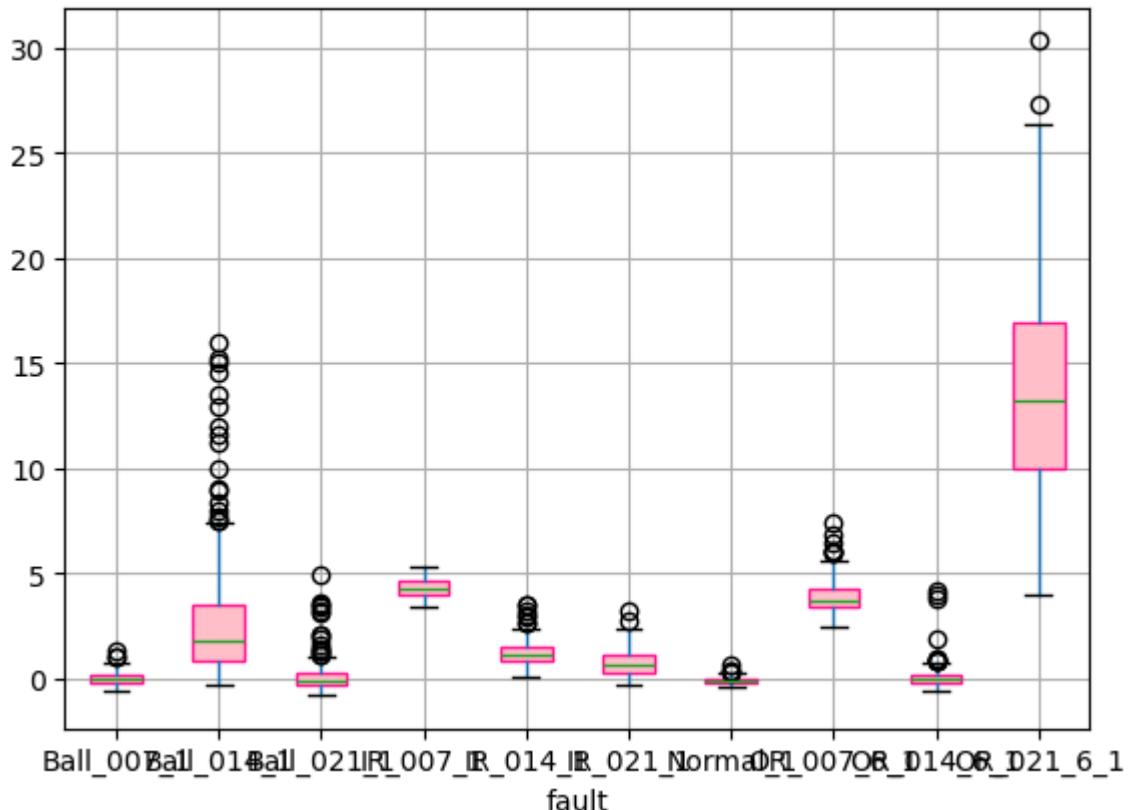
for f in features:
    plt.figure(figsize=(10,8))
    df.boxplot(column=f, by='fault', patch_artist=True,
               boxprops=dict(facecolor='pink', color='deeppink'))
    df.boxplot(column=f, by='fault')
    plt.title(f'{f} Distribution by Fault Class')
    plt.suptitle('')
    plt.xlabel('Fault')
    plt.ylabel(f)
    plt.show()
```

<Figure size 1000x800 with 0 Axes>

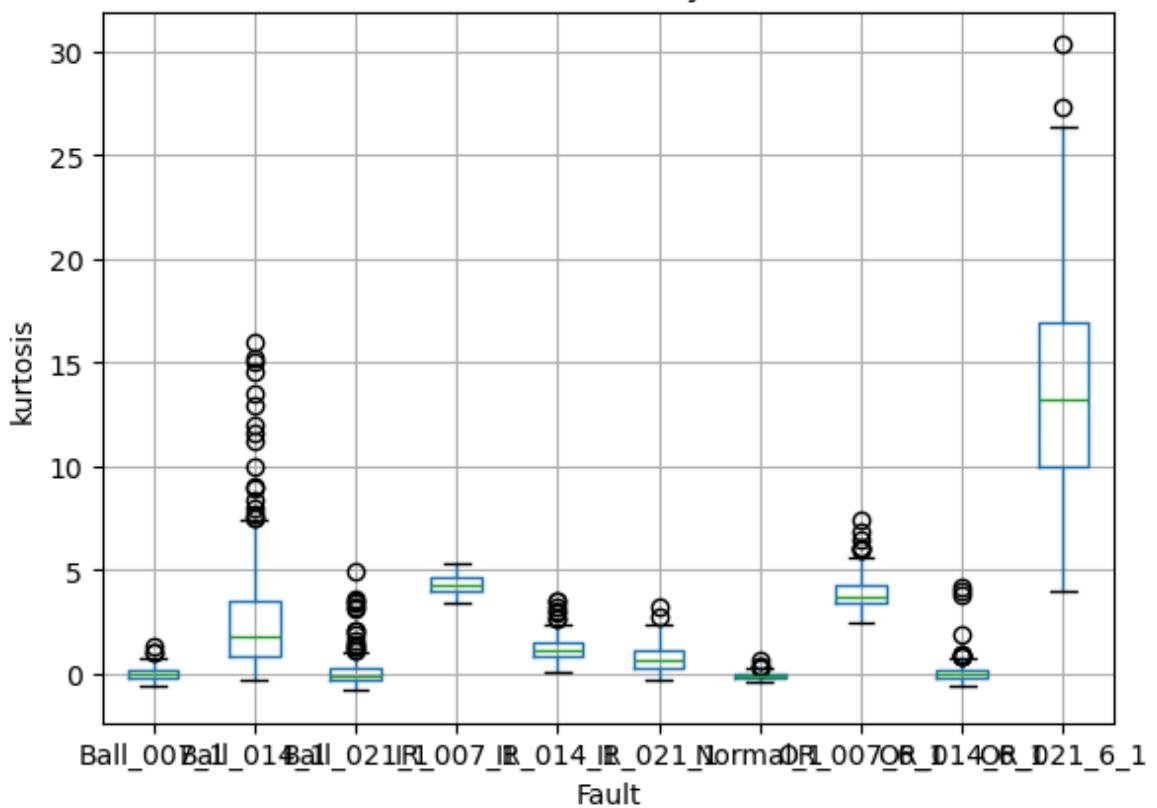


<Figure size 1000x800 with 0 Axes>

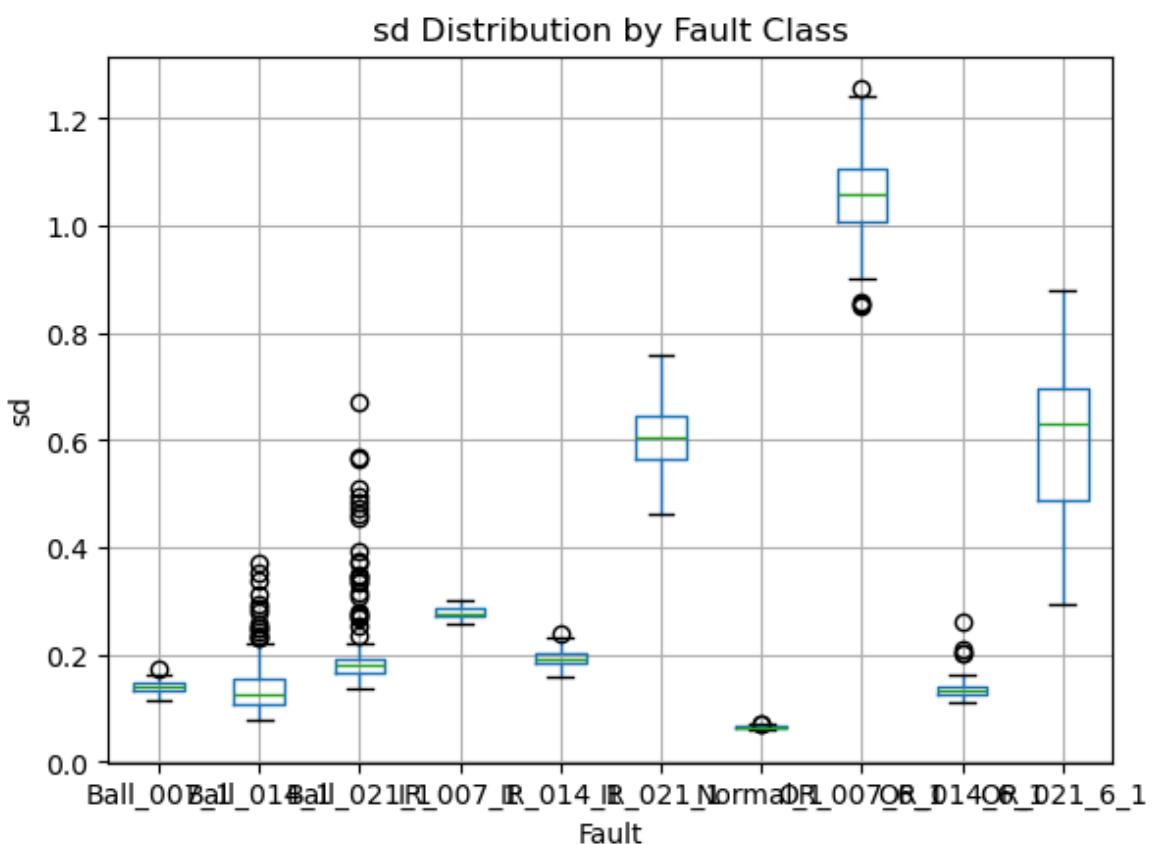
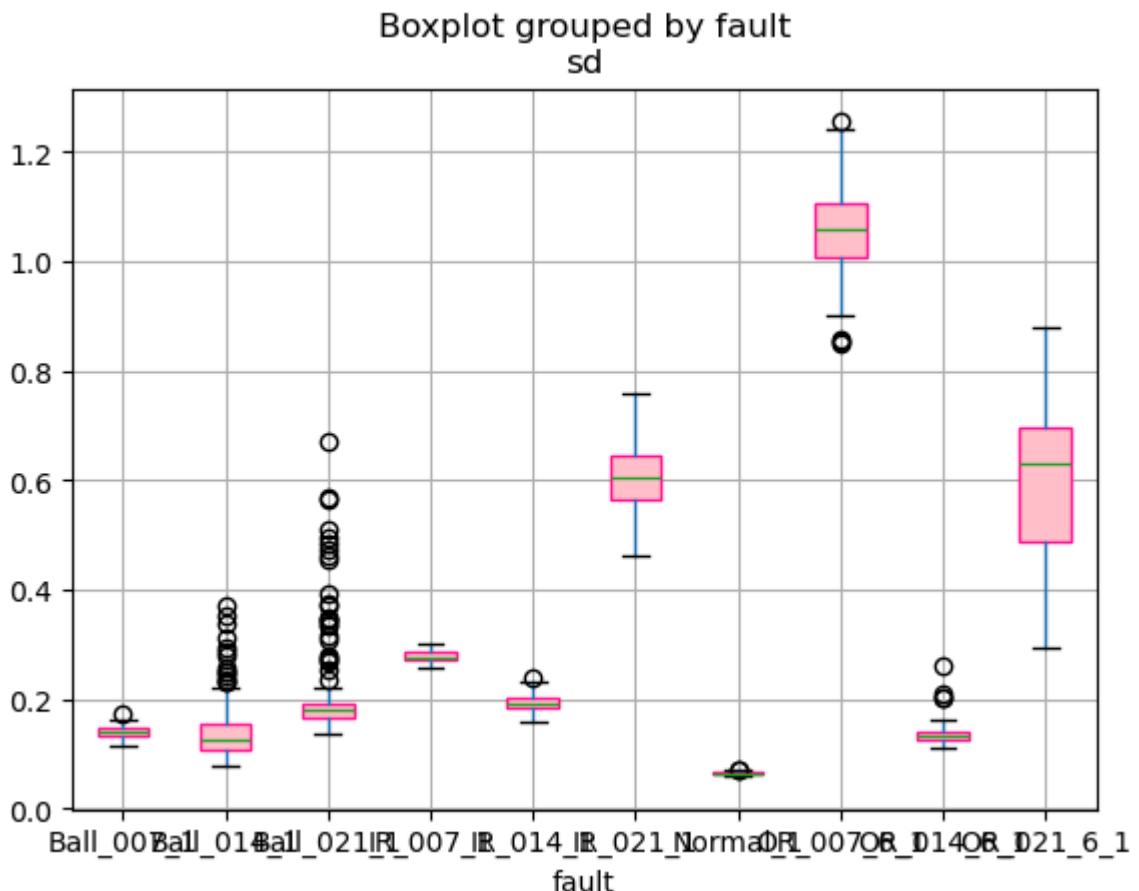
Boxplot grouped by fault kurtosis



kurtosis Distribution by Fault Class



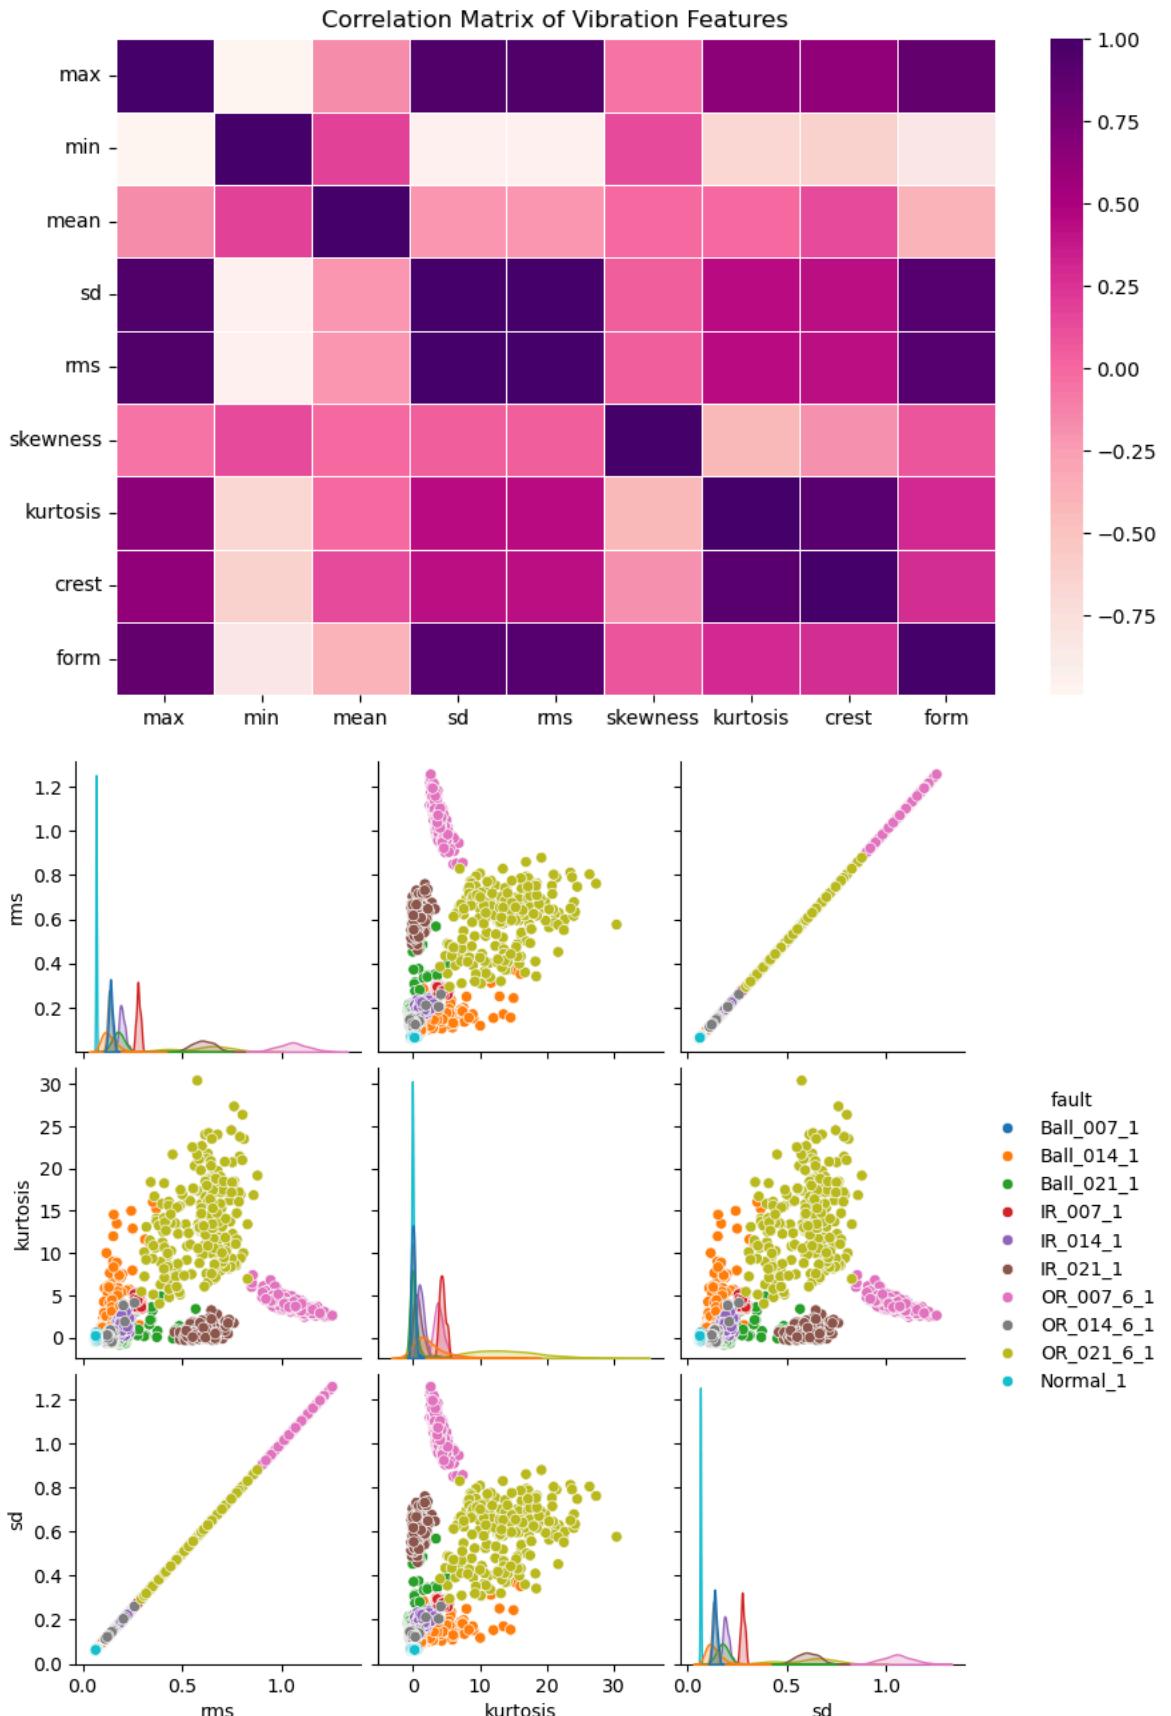
<Figure size 1000x800 with 0 Axes>



In [56]: `import seaborn as sns`

```
plt.figure(figsize=(10,6))
sns.heatmap(X.corr(), cmap='RdPu', linewidths=0.5)
plt.title('Correlation Matrix of Vibration Features')
plt.show()
```

```
sns.pairplot(df[['rms', 'kurtosis', 'sd', 'fault']], hue='fault')
plt.show()
```



```
In [59]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
```

```

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

X = df.drop(columns=['fault'])
y = df['fault']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

log_reg = LogisticRegression(max_iter=1000)

log_reg.fit(X_train_scaled, y_train)

y_pred_lr = log_reg.predict(X_test_scaled)

print("Logistic Regression Accuracy:",
      accuracy_score(y_test, y_pred_lr))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_lr))

cm_lr = confusion_matrix(y_test, y_pred_lr)

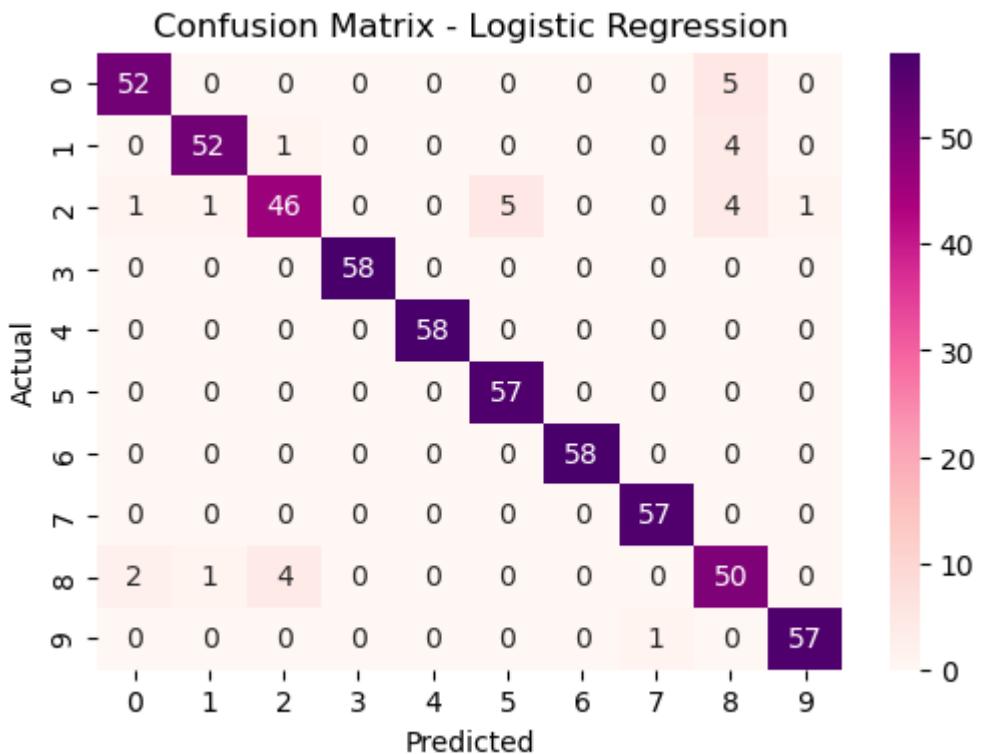
plt.figure(figsize=(6,4))
sns.heatmap(cm_rf, annot=True, fmt='d', cmap='RdPu')
plt.title('Confusion Matrix - Logistic Regression')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

Logistic Regression Accuracy: 0.9130434782608695

Classification Report:

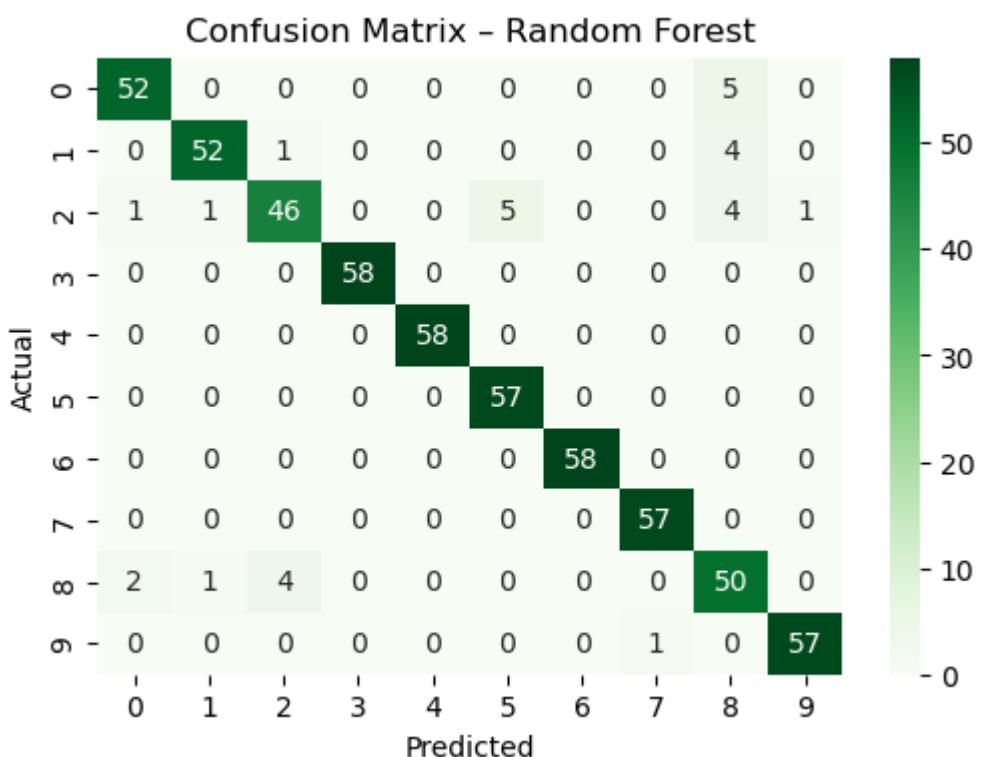
| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Ball_007_1 | 0.92 | 0.96 | 0.94 | 57 |
| Ball_014_1 | 0.91 | 0.74 | 0.82 | 57 |
| Ball_021_1 | 0.80 | 0.71 | 0.75 | 58 |
| IR_007_1 | 0.97 | 1.00 | 0.98 | 58 |
| IR_014_1 | 1.00 | 1.00 | 1.00 | 58 |
| IR_021_1 | 0.90 | 1.00 | 0.95 | 57 |
| Normal_1 | 0.93 | 0.98 | 0.96 | 58 |
| OR_007_6_1 | 0.98 | 1.00 | 0.99 | 57 |
| OR_014_6_1 | 0.73 | 0.79 | 0.76 | 57 |
| OR_021_6_1 | 0.98 | 0.95 | 0.96 | 58 |
| accuracy | | | 0.91 | 575 |
| macro avg | 0.91 | 0.91 | 0.91 | 575 |
| weighted avg | 0.91 | 0.91 | 0.91 | 575 |



```
In [85]: cm_rf = confusion_matrix(y_test, y_pred_rf)

plt.figure(figsize=(6,4))
sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Greens')
plt.title('Confusion Matrix – Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

importances = rf.feature_importances_
features = X.columns
```

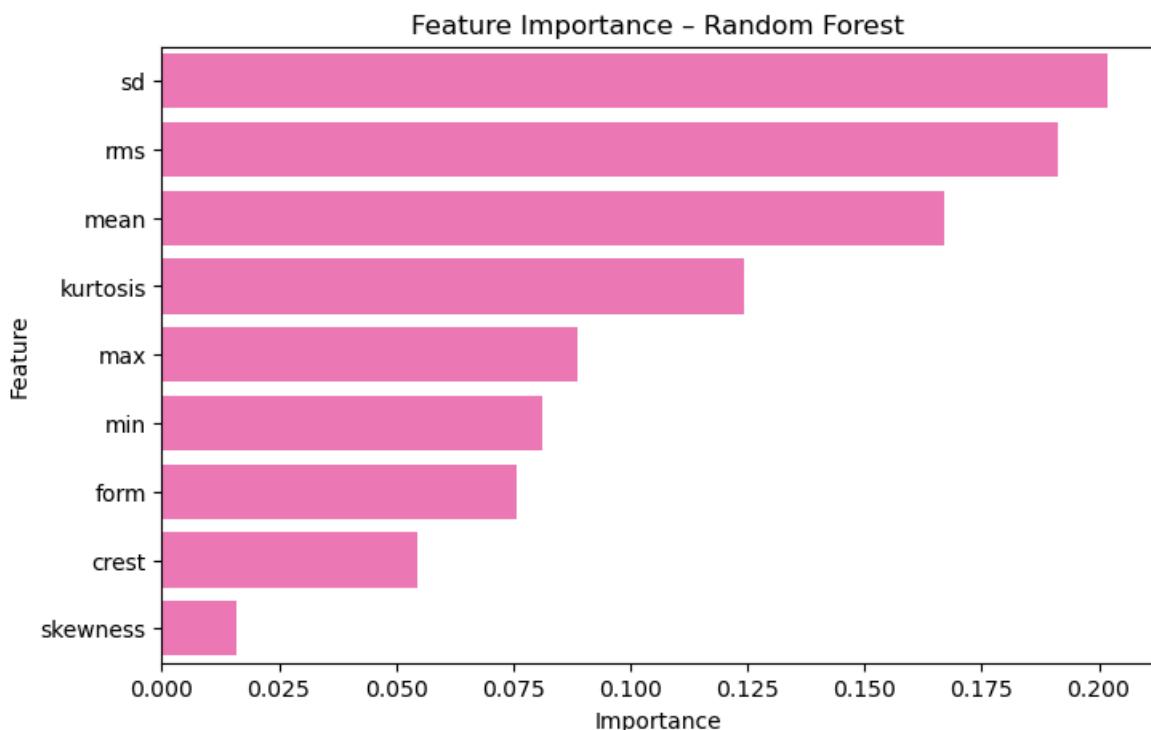


```
In [61]: importances = rf.feature_importances_
features = X.columns

importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(8,5))
sns.barplot(x='Importance', y='Feature', data=importance_df, color='hotpink')
plt.title('Feature Importance – Random Forest')
plt.show()

importance_df
```



```
Out[61]: Feature    Importance
      3        sd    0.201836
      4        rms    0.191224
      2       mean    0.166934
      6     kurtosis    0.124166
      0       max    0.088654
      1       min    0.081254
      8       form    0.075587
      7      crest    0.054461
      5   skewness    0.015883
```

```
In [84]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

X = df.drop(columns=['fault'])
y = df['fault']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=42, stratify=y
)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

log_reg = LogisticRegression(max_iter=1000)

log_reg.fit(X_train_scaled, y_train)

y_pred_lr = log_reg.predict(X_test_scaled)

print("Logistic Regression Accuracy:",
      accuracy_score(y_test, y_pred_lr))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_lr))

cm_lr = confusion_matrix(y_test, y_pred_lr)

plt.figure(figsize=(6,4))
sns.heatmap(cm_lr, annot=True, fmt='d', cmap='RdPu')
plt.title('Confusion Matrix - Logistic Regression')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

cm_rf = confusion_matrix(y_test, y_pred_rf)

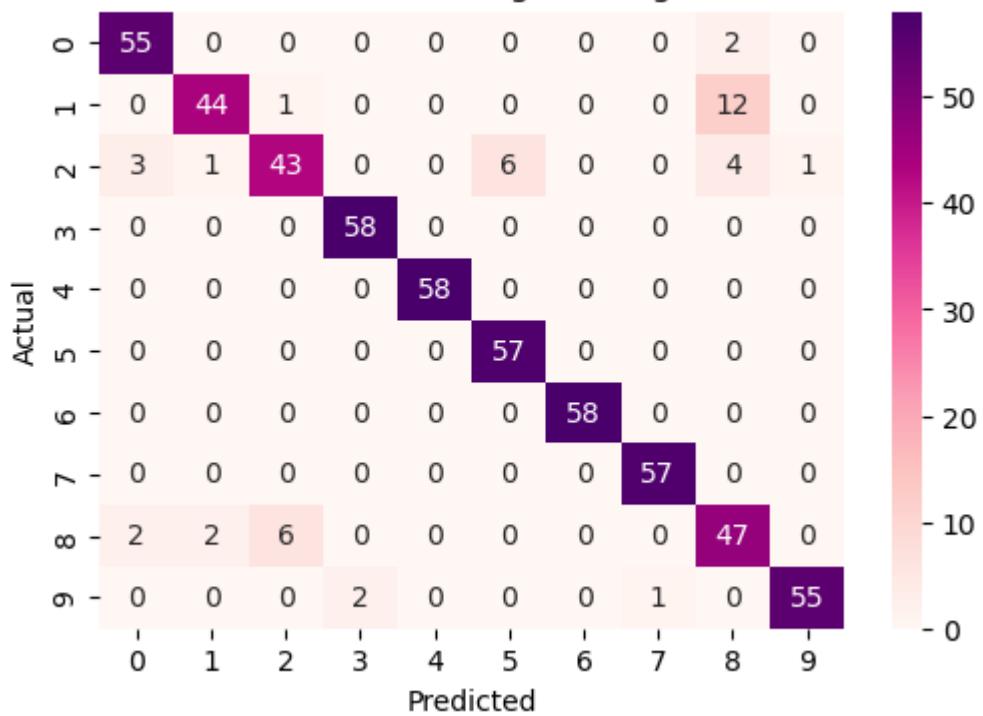
plt.figure(figsize=(6,4))
sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Greens')
plt.title('Confusion Matrix - Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

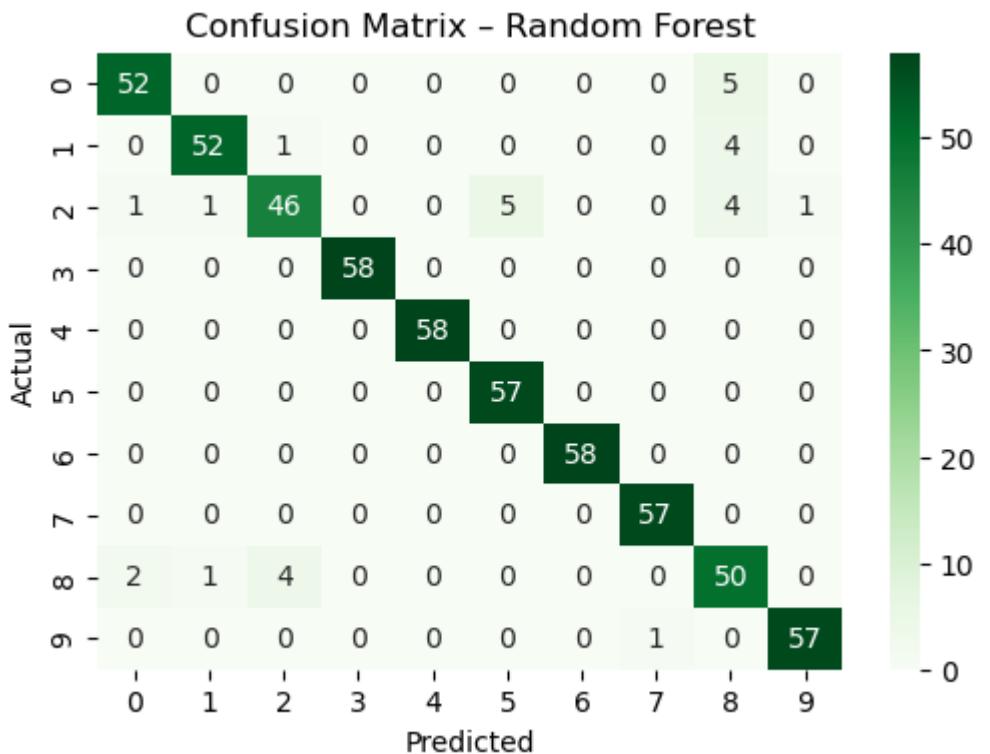
Logistic Regression Accuracy: 0.9252173913043479

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Ball_007_1 | 0.92 | 0.96 | 0.94 | 57 |
| Ball_014_1 | 0.94 | 0.77 | 0.85 | 57 |
| Ball_021_1 | 0.86 | 0.74 | 0.80 | 58 |
| IR_007_1 | 0.97 | 1.00 | 0.98 | 58 |
| IR_014_1 | 1.00 | 1.00 | 1.00 | 58 |
| IR_021_1 | 0.90 | 1.00 | 0.95 | 57 |
| Normal_1 | 1.00 | 1.00 | 1.00 | 58 |
| OR_007_6_1 | 0.98 | 1.00 | 0.99 | 57 |
| OR_014_6_1 | 0.72 | 0.82 | 0.77 | 57 |
| OR_021_6_1 | 0.98 | 0.95 | 0.96 | 58 |
| accuracy | | | 0.93 | 575 |
| macro avg | 0.93 | 0.93 | 0.92 | 575 |
| weighted avg | 0.93 | 0.93 | 0.92 | 575 |

Confusion Matrix - Logistic Regression





```
In [87]: #is the machine faulty or not?
#is the machine faulty or not?
df['binary_fault'] = df['fault'].apply(
    lambda x: 0 if x == 'Normal_1' else 1
)

# 2. Separate features and the new binary target
# [cite: 733, 734]
X_bin = df.drop(columns=['fault', 'binary_fault'])
y_bin = df['binary_fault']

# 3. Fix the train_test_split syntax
# [cite: 735-737]
from sklearn.model_selection import train_test_split

Xb_train, Xb_test, yb_train, yb_test = train_test_split(
    X_bin,
    y_bin,
    test_size=0.25,
    random_state=42,
    stratify=y_bin
)

# 4. Train and Predict
# [cite: 738, 739]
from sklearn.ensemble import RandomForestClassifier
rf_bin = RandomForestClassifier(n_estimators=200, random_state=42)
rf_bin.fit(Xb_train, yb_train)
yb_pred = rf_bin.predict(Xb_test)

# 5. Generate Corrected Confusion Matrix
# [cite: 740-754]
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```

cm_bin = confusion_matrix(yb_test, yb_pred, labels=[0, 1])

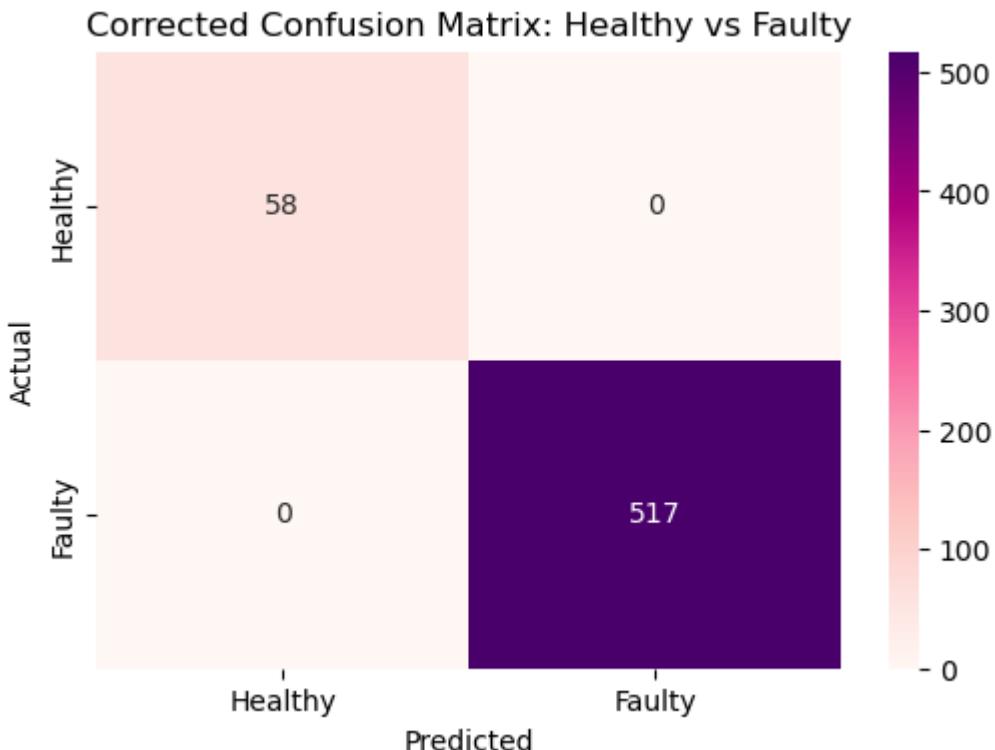
plt.figure(figsize=(6,4))
sns.heatmap(
    cm_bin,
    annot=True,
    fmt='d',
    cmap='RdPu',
    xticklabels=['Healthy', 'Faulty'],
    yticklabels=['Healthy', 'Faulty']
)
plt.title('Corrected Confusion Matrix: Healthy vs Faulty')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Changed from best_rf to rf_bin since that's the model we trained
importances = rf_bin.feature_importances_

# Changed X to X_bin to match the features used in training
importance_df = pd.DataFrame({
    'Feature': X_bin.columns,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

importance_df

```



Out[87]:

| | Feature | Importance |
|---|----------|------------|
| 3 | sd | 0.310511 |
| 4 | rms | 0.288796 |
| 0 | max | 0.186632 |
| 1 | min | 0.172125 |
| 8 | form | 0.035206 |
| 5 | skewness | 0.002418 |
| 2 | mean | 0.001989 |
| 6 | kurtosis | 0.001279 |
| 7 | crest | 0.001045 |

In [86]: *#assuming these values to show the benefits of the model*

```
cost_per_failure = 50000      # ₹ (downtime + repair)
annual_failures = 12          # failures per machine per year
downtime_reduction = 0.4      # 40% reduction using PdM

annual_loss = cost_per_failure * annual_failures
annual_savings = annual_loss * downtime_reduction

annual_loss, annual_savings
```

Out[86]: (600000, 240000.0)